

ASSIGNMENT – LAB 1.4 : Log Analysis & Spark UI Deep Dive

Learning Outcomes

By completing this assignment, you should be able to:

- Generate synthetic **web server logs** in a realistic text format.
- Parse semi-structured logs using **regex** into structured columns.
- Perform log-based **status/error**, **performance**, and **traffic** analysis with Spark.
- Persist cleaned logs and summaries as **Parquet/CSV**.
- Use the **Spark UI**(Jobs, Stages, Storage, Environment, Executors) to:
 - Read DAG visualizations and stages.
 - Understand shuffles, caching, and resource usage.
 - Relate Spark code to physical execution.

0. Context & Prerequisites

You are now instrumenting a simulated e-commerce website from an **operational analytics** viewpoint.

Prerequisites:

- `spark-data/ecommerce/` already exists (from previous labs).
- Spark environment working:
 - Either `master="local[*]"` or your cluster (`spark://localhost:7077`).
 - Spark Application UI accessible at `http://localhost:4040` while jobs are running.

Your goal: go from **raw log lines** → **structured analysis** → **UI interpretation**.

1. Part A – Generate Web Server Logs

Goal: Create a realistic text log file for a web server.

Task A.1 – Implement `generate_logs.py`

From your **project root** (e.g. `data-engineering-course`), create:

`generate_logs.py`

The script must:

1. Generate **10,000** synthetic log lines.

2. Use random combinations of:

- IPv4 addresses,
- HTTP methods (`GET`, `POST`, etc.),
- URLs (`/home`, `/products`, `/cart`, etc.),
- HTTP status codes (200, 304, 404, 500, 503),
- response times (latency in ms),
- user agents (desktop, mobile, etc.).

3. Emit log lines in an Apache-like format:

```
IP -- [timestamp] "METHOD ENDPOINT HTTP/1.1" STATUS RESPONSETIME "-"
```

4. Write all lines to:

```
spark-data/ecommerce/web_logs.txt
```

Then run:

```
python generate_logs.py
```

Check:

- File exists: `spark-data/ecommerce/web_logs.txt`.
- Script reports **10,000** generated log entries.

2. Part B – Parse and Analyze Logs with Spark

Goal: Parse the raw logs into structured columns and perform analytics.

Task B.1 – Create `lab4_log_analysis.py`

From your project root, create:

```
lab4_log_analysis.py
```

The script must:

1. Create a SparkSession

- App name: `"Day1-LogAnalysis"`.
- Master: `local[*]` (or your cluster master).
- `spark.driver.memory = "2g"`.

2. Load raw logs

- Read from `"spark-data/ecommerce/web_logs.txt"` using
`spark.read.text(...)`.
- Print:
 - Number of log lines (`raw_logs.count()`).
 - A small sample with `raw_logs.show(3, truncate=False)`.

3. Parse the log format using regex

- Define a pattern similar to:

```
log_pattern = r'(\S+) - - \[(\w+:/]+\s[+\-]\d{4})\] (\S+) (\S+)
```

- Extract:
 - `ip`
 - `timestamp` (string)
 - `method`
 - `endpoint`
 - `protocol`
 - `status` (cast to `int`)
 - `response_time_ms` (cast to `int`)

using `regexp_extract` on column "value".

- Build a `parsed_logs` DataFrame with these columns.
- Print its schema and the first 10 rows.

4. Data quality check

- Compute:
 - `total_logs = parsed_logs.count()`.
 - `valid_logs = parsed_logs.filter(col("ip") != "").count()`.
 - `invalid_logs = total_logs - valid_logs`.
- Print these three metrics.
- Create `logs = parsed_logs.filter(col("ip") != "")` as the cleaned DataFrame.

5. Basic analytics

- Status code distribution:

```
logs.groupBy("status").count().orderBy(desc("count"))
```

- HTTP method distribution:

```
logs.groupBy("method").count().orderBy(desc("count"))
```

- Top 10 most visited pages (by `endpoint`).

6. Error analysis

- Count total `4xx` and `5xx` errors.
- Top 404 pages:

```
logs.filter(col("status") == 404) \
    .groupBy("endpoint") \
    .count() \
    .orderBy(desc("count"))
```

- Pages causing `5xx` errors (group by `endpoint` and `status`).

7. Performance analysis

- Global response time statistics:
 - count, min, max, avg for `response_time_ms` .
- Slowest endpoints (by average latency), for endpoints with >10 requests.
- Top 10 slowest individual requests (highest `response_time_ms`).

8. Traffic patterns

- Add `parsed_timestamp` from string `timestamp` :

```
logs_with_time = logs.withColumn(  
    "parsed_timestamp",  
    to_timestamp(col("timestamp"), "dd/MMM/yyyy:HH:mm:ss Z"),  
)
```

- Compute **traffic by hour**:

- Add `hour(parsed_timestamp)` as `hour` .
- Group by `hour`, count, order by `hour` .

9. User behavior

- Top 10 most active IPs (by request count).
- Compute `requests_per_ip = logs.groupBy("ip").count()` and show `describe()` .
- Potential bots: IPs with `count > 100` , ordered by count.

10. Save processed data and summary

- Save `logs_with_time` as Parquet to:

```
spark-data/ecommerce/processed_logs
```

```
using mode ("overwrite") .
```

- Build a `summary` DataFrame with metrics such as:
 - Total Requests
 - Valid Requests
 - Unique IPs

- Unique Pages
- 4xx Errors
- 5xx Errors

- Show the summary, then write it as a single CSV (use `coalesce(1)`) to:

```
spark-data/ecommerce/log_summary
```

11. **Stop Spark** with `spark.stop()`.
-

3. Part C – Spark UI Exploration

Goal: Learn to read the Spark UI and connect it to your code.

Suggestion:

1. Open `http://localhost:4040` in your browser.

2. Then run:

```
python lab4_log_analysis.py
```

3. Refresh the UI while the job runs.

Task C.1 – Jobs Tab

- Go to: `http://localhost:4040/jobs/`.
- Count how many jobs were executed.
- Identify the **longest-running job** (by duration).
- For that job:
 - How many stages does it have?
 - Take a screenshot of the **DAG Visualization**.

Write short answers:

- Why does this job have multiple stages?
 - Which operations in your code likely caused shuffles (and therefore extra stages)?
-

Task C.2 – Stage Details

- Click into one **stage** of the longest job.
- Observe:
 - Number of tasks.
 - Shuffle read/write size.
 - Task durations.

Answer:

- Are task durations similar or is there skew?
 - How much data was shuffled in this stage?
-

Task C.3 – Storage Tab & Caching Demo

Create and run a small script: `lab4_caching_demo.py` that:

- Loads `web_logs.txt` into a DataFrame `logs`.
- **Without caching:**
 - Run `logs.count()` twice, measure and print both times.
- **With caching:**
 - Create `logs_cached = logs.cache()`.
 - Run `logs_cached.count()` twice, measure times.
 - Print a computed speedup.

Add an `input("Press Enter to continue...")` before stopping Spark so the UI stays visible.

While it's paused:

- Go to `http://localhost:4040/storage/` and observe:
 - Cached dataset name.
 - Number of cached partitions.
 - Storage level (e.g., MEMORY_ONLY).
 - Memory usage.

Answer:

- Is the second `count()` noticeably faster when cached?

- What does the Storage tab tell you about how Spark keeps data in memory?
-

Task C.4 – Environment Tab

Navigate to <http://localhost:4040/environment/> and inspect:

- `spark.master`
- `spark.driver.memory`
- Java version

Answer briefly:

- What master is your app connected to?
 - How much memory is allocated to the driver?
-

Task C.5 – Executors Tab

Navigate to <http://localhost:4040/executors/>.

Observe:

- Number of executors (in local mode you typically see 1 executor + driver).
- Executor memory, active tasks, shuffle metrics, and GC time.

Answer:

- How many executors (excluding driver)?
 - Is GC time significant compared to task time?
-

4. Part D – Challenge Problems (Advanced, for extra credit)

Using `logs_with_time` from `lab4_log_analysis.py`, implement the following at the bottom of the script (or in a separate `lab4_challenges.py`), and print the results.

Challenge 1 – Bounce Rate

| **Definition:** percentage of unique IPs that made **exactly one request**.

Steps:

- Group `logs` (or `logs_with_time`) by `ip` and count requests.
- Identify IPs with `count == 1`.
- Compute:

```
bounce_rate = bounced_ips / total_unique_ips * 100
```

Print bounce rate in percentage.

Challenge 2 – Conversion Funnel

We define a simple funnel:

```
/home → /products → /product/* → /cart → /checkout
```

Goal: estimate how many visitors (unique IPs) reach each step.

Hints:

- Define boolean flags or filtered DataFrames for each step:
 - `/home`
 - `/products`
 - endpoints starting with `/product/`
 - `/cart`
 - `/checkout`
 - For each step, compute the set or count of **unique IPs** that visited at least one matching endpoint.
 - Compare counts across steps to estimate drop-off.
-

Challenge 3 – Peak Traffic Hour

- Which hour of the day has the most requests?
- How does the traffic distribute across the 24 hours?

Steps:

- Use `hour(parsed_timestamp)` from `logs_with_time`.
- Group by hour and count.
- Order by hour and print table.

- compute percentages.
-

Challenge 4 – Suspicious Activity

Detect potentially abusive patterns, for example:

- IPs with **>50 requests per minute**.
- IPs with a **high proportion of 404s** (e.g. >50% of their requests).

Hints :

- Use `groupBy(ip, window(parsed_timestamp, "1 minute"))` to compute requests per minute.
 - For 404 rate: group by `ip`, compute total requests and 404 requests, then compute ratio.
-

Challenge 5 – Performance Bottlenecks

- List endpoints with **average response_time_ms > 2000 ms**, for endpoints with more than a small number of hits (e.g. >10).
- For these endpoints, analyze relation with status code:
 - Group by `endpoint, status` and compute `avg(response_time_ms)` .

BONUS: Analyze whether some hours have systematically higher average response times.

5. Hints (High Level)

- For regex parsing: always inspect 3-5 sample lines before trusting the pattern.
 - For time parsing: ensure timestamp format in `to_timestamp` is exactly `"dd/MMM/yyyy:HH:mm:ss Z"` .
 - For “per IP” metrics: **groupBy("ip")** first, then aggregate.
 - For “> X requests” conditions: use `filter(col("count") > X)` on an aggregated DataFrame.
 - For the funnel: unique IP counts per step are enough (no need for full session tracking in this lab).
-

6. Common Pitfalls

- Wrong log path: `spark-data/ecommerce/web_logs.txt` must exist before running analysis.
 - Regex not matching because of typos or spaces → leads to empty columns (check sample output).
 - Forgetting to cast `status` and `response_time_ms` to `int` → numeric aggregations fail or behave strangely.
 - Calling `count()` on `raw_logs` and `parsed_logs` without understanding extra jobs (normal; each action is a job).
 - Trying to compute “per-minute” windows without converting to timestamp.
-

7. Deliverables

You should submit:

1. Code:

- `generate_logs.py`
- `lab4_log_analysis.py`
- `lab4_caching_demo.py`
- `lab4_challenges.py`

2. Outputs / Evidence:

- Console output (or text export) from `lab4_log_analysis.py`.
- Console output from `lab4_caching_demo.py` showing timing before/after caching and computed speedup.

3. Screenshots of Spark UI:

- Jobs tab (with at least one DAG visualization).
- A Stage detail page (for a chosen job).
- Storage tab showing the cached dataset from `lab4_caching_demo.py`.
- Executors tab.

4. Short written answers:

- Answers to the UI questions in Part C (Jobs, Stages, Storage, Environment, Executors).
- If challenges attempted:
 - Results + 2–3 lines explaining each (bounce rate, peak hour, suspicious IPs, etc.).